

Hedge Fund Panel Forecasting

Multi-Horizon Time-Series Modelling under a Weighted Skill Score Metric

Author 1

Author 2

Author 3

Author 4

DA-IICT, Semester 6
Applied Forecasting Methods, 2026

Abstract

Abstract to be written after final results have landed. This project studies a panel of approximately 37,000 anonymised time series taken from a Kaggle hedge-fund forecasting competition. The target is forecast at four lookahead horizons under a weighted skill score that concentrates roughly two thirds of the evaluation mass in the top one percent of rows. The first phase of the work, presented in this draft, establishes the structure of the panel, the diagnostic conditions for classical and machine-learning models, and three findings that were not surfaced in the midway report: the four horizons are coupled cumulations of a single short-step process, the maximum panel weight lives at horizon three, and per-horizon target scaling differs by a factor of more than four.

Contents

1	Problem Statement	3
1.1	What is being forecast	3
1.2	Why the problem is hard	3
1.3	Why the problem is interesting	3
1.4	Scope of this report	4
2	Forecast Targets and Series Identity	5
2.1	Panel keys	5
2.2	The four horizons are coupled	5
2.3	Chronological constraint	5
2.4	Eligibility for per-series studies	6
3	The Weighted Skill Score	7
3.1	Definition	7
3.2	Intuitive meaning	7
3.3	Practical consequences for this project	7
4	Dataset and Exploratory Analysis	9
4.1	Schema, panel scale, and train–test comparability	9
4.2	Missing values	9
4.3	Target distribution	10
4.4	Weight distribution and metric concentration	10
4.5	Series coverage and split feasibility	11
4.6	Representative-series selection	11
4.7	Stationarity diagnostics	12
4.8	Visual stationarity — rolling statistics	13
4.9	Lag structure — ACF and PACF	14
4.10	STL decomposition	14
4.11	Feature audit	14
4.12	Cross-series relationships	15
4.13	Multi-horizon coupling	15
4.14	Cumulation hypothesis	16
4.15	Per-horizon weight dominance	17
4.16	EDA-driven modelling implications	18
5	Classical Baselines	21
5.1	Setup and the four representative series	21
5.2	Model lineup	21
5.3	Evaluation protocol	22
5.4	Per-series leaderboard	23
5.5	Skip-status under adaptive splits	23
5.6	Forecast overlays	23
5.7	What this study can and cannot say	24

Section 1. Problem Statement

We work with a publicly released Kaggle competition dataset titled *Hedge fund – time-series forecasting*. The competition sponsor releases a panel of anonymised numerical series with engineered features and asks competitors to forecast a continuous target at four lookahead horizons. The dataset is intentionally anonymised: identifiers, feature names, and the underlying financial domain are stripped. What remains is a clean, mathematically well-defined supervised forecasting problem with a non-trivial weighting scheme.

1.1 What is being forecast

Each row of the panel is identified by the tuple

`(code, sub_code, sub_category, horizon, ts_index)`

and carries a continuous target y_{target} , a non-negative competition weight, and 86 anonymised numerical covariates of the form `feature_a`, `feature_b`, ... The integer time index `ts_index` runs over the contiguous range 1–3601 in the public training partition and 3602–4376 in the held-out test partition. The training panel contains 5,337,414 rows and the test panel contains 1,447,107 rows.

The four forecast horizons are $H \in \{1, 3, 10, 25\}$, interpreted as the lookahead in time units. For each `(code, sub_code, sub_category)` the panel typically carries all four horizons at the same `ts_index`, so the natural unit of analysis is the 4-tuple key above. There are 36,923 distinct series in train under that definition.

1.2 Why the problem is hard

The dataset combines three well-known difficulties that make naive baselines a bad guide:

1. **Heavy-tailed target.** The target distribution is centred near zero with an interquartile range of approximately ± 0.05 , but the full range runs from roughly -2200 to $+2300$. The 1% and 99% percentiles are at approximately ± 82 . Plain mean-squared losses are therefore extremely sensitive to a small number of tail rows.
2. **Concentrated evaluation weight.** The competition metric is sample-weighted, and the weights are extraordinarily skewed: the median weight is approximately 1.7×10^3 while the maximum is 1.4×10^{13} . As a result, the top one percent of rows by weight carry approximately 64% of all evaluation mass, and the top five percent carry approximately 92%. Any unweighted analysis can recommend the wrong model.
3. **Limited per-series history at long horizons.** Although the panel is large in aggregate, only about 5% of individual series cleanly cross the validation cutoff used in this work (`ts_index` = 2880). Some series have as few as seven training points before the cutoff, which makes per-series classical fitting infeasible for many model families.

1.3 Why the problem is interesting

Beyond being a real competitive forecasting problem, three properties make the dataset a useful subject for a methodological course project:

- *Anonymisation forces statistical reasoning.* Without semantic feature names we cannot fall back on domain priors. Every claim must be supported by a measurable property of the data.
- *The weighting scheme rewards precision on rare rows.* A model that improves average behaviour while degrading on the high-weight rows can lose to a simpler model. This is a mechanically different objective from standard regression and warrants explicit study.

- *The four horizons are coupled.* As we show in Section 4.14, the targets at $H \in \{3, 10, 25\}$ are well approximated by rolling cumulative sums of the $H = 1$ process. This structure is not advertised in the dataset description and admits a modelling strategy that the midway report did not consider.

1.4 Scope of this report

The remainder of this report covers the diagnostic and exploratory phase of the project. Section 2 formalises the panel keys and the chronological constraint that any validation strategy must respect. Section 3 develops the official competition metric in detail, with worked numerical examples and a comparison against standard error metrics. Section 4 carries out a full exploratory analysis: schema sanity, missingness, target and weight distributions, stationarity diagnostics, autocorrelation structure, feature audit, and three findings about the multi-horizon structure of the panel that constrain every modelling decision in subsequent chapters of the project. Modelling chapters (classical baselines, global gradient-boosted models, the multi-horizon aggregation experiment, probabilistic forecasting, and ensembles) will be added as those notebooks land.

Section 2. Forecast Targets and Series Identity

A panel forecasting problem only makes sense once the unit of analysis is unambiguous. This section fixes the definitions used throughout the project.

2.1 Panel keys

A row of the panel has primary keys (`code`, `sub_code`, `sub_category`, `horizon`, `ts_index`). We treat the 4-tuple

$$s \equiv (\text{code}, \text{sub_code}, \text{sub_category}, \text{horizon})$$

as the natural *series identifier*, and treat `ts_index` as the time axis. There are 23 codes, 180 distinct sub-codes, 5 sub-categories, and 4 horizons in train, yielding 36,923 distinct series. Test uses a narrower sub-code support of 47 distinct values, which is a coverage shift rather than a schema mismatch (every test category level already exists in train).

2.2 The four horizons are coupled

A naive reading of the schema would treat the four horizons of the same (`code`, `sub_code`, `sub_category`) as four independent series, simply distinguished by an extra integer label. Two empirical facts in Section 4 contradict this view:

1. For approximately 1.21 million distinct (`code`, `sub_code`, `sub_category`, `ts_index`) tuples, all four horizons are simultaneously present (Section 4.13).
2. For the same triple, the target at horizon k is closely approximated by the rolling sum of the $H = 1$ target over the next k steps:

$$y_t^{H_k} \approx \sum_{i=0}^{k-1} y_{t+i}^{H_1}.$$

We verify this on a stratified sample of 500 triples in Section 4.14, and find median Pearson correlations of 0.96, 0.94 and 0.91 at $k = 3, 10, 25$ respectively.

The practical consequence is that fitting four independent regressors per triple discards a strong constraint that the data construction itself imposes. We exploit this constraint in the H1-aggregation experiment scheduled for a later notebook.

2.3 Chronological constraint

The training partition runs over `ts_index` $\in [1, 3601]$ and the test partition over $[3602, 4376]$. Because test is a strict temporal suffix of train, any internal validation strategy must also be chronological, with all training points strictly preceding the validation points in time. Random or leave-one-series-out cross-validation would mix future information into earlier states and overstate generalisation.

We use a single canonical chronological partition throughout the project, fixed in the shared utilities at `cf.splits`:

- **Train:** `ts_index` ≤ 2880 .
- **Meta** (used only for fitting ensemble weights): $2880 < \text{ts_index} \leq 3240$.
- **Holdout** (final scored window): $3240 < \text{ts_index} \leq 3601$.
- **Test** (`ts_index` ≥ 3602): unlabeled, used only for submission-style sanity checks.

The cutoff `ts_index` = 2880 matches the cutoff used in the midway report, preserving comparability with prior diagnostic figures. The meta/holdout boundary at 3240 places approximately one third of the post-cutoff window into the meta slice, leaving the remaining two thirds for final scoring.

2.4 Eligibility for per-series studies

Some classical model families (ARIMA, ARMA, MA, AR with order ≥ 2) require a minimum training history. For per-series diagnostic studies we restrict attention to series that both cross the train/validation cutoff and have at least 120 training rows. There are 1,930 such *eligible* series in the panel, roughly 5% of the 36,923 total. Global pooled models and the H1-aggregation experiment use every row whose `ts_index` falls inside the relevant split, so they are not constrained by per-series length.

Section 3. The Weighted Skill Score

The competition is decided by a single scalar: the *weighted skill score*. Because the rest of the project uses this metric for every model comparison, ensemble weight, and bootstrap interval, it deserves a dedicated section. We state the formula, give the intuitive reading, and note the two practical consequences that carry through every modelling decision.

3.1 Definition

Given truths y_i , predictions $\hat{y}_i$, and non-negative weights w_i for $i = 1, \dots, n$, the weighted skill score is

$$\text{score}(y, \hat{y}; w) = \sqrt{1 - \text{clip}_{[0,1]} \left(\frac{\sum_i w_i (y_i - \hat{y}_i)^2}{\sum_i w_i y_i^2} \right)} \quad (1)$$

where $\text{clip}_{[0,1]}(x) = \min(\max(x, 0), 1)$. The function returns 1 for a perfect prediction and 0 when the prediction is no better than the all-zero forecast.

3.2 Intuitive meaning

Strip away the symbols and the score answers a single question: *how much of the weighted target “energy” is left unexplained by the prediction?*

The fraction inside the square root, $\sum w_i (y_i - \hat{y}_i)^2 / \sum w_i y_i^2$, is exactly that — the unexplained share of weighted squared truth. It is zero when the prediction is perfect and one when the prediction is identically zero. The clip caps it at one so that no model can score worse than the trivial all-zero forecast, and the outer square root puts the result on a familiar correlation-like scale where 1 is excellent, 0 is no skill, and intermediate values read roughly the way a Pearson correlation of the same magnitude reads.

A few useful reference points:

- **Skill** = 1 — perfect predictions on every weighted row.
- **Skill** = 0 — the model is no better than predicting zero. This is the natural “did you learn anything” threshold.
- **Skill** $\approx$ 0.5 — roughly three quarters of the weighted target energy is still unexplained. In feel, this is comparable to a Pearson correlation of 0.5: a real but modest signal.
- **Skill** $\approx$ 0.95 — only ten percent of energy unexplained. Strong, near the practical ceiling for a competitive entry.

The reason the metric is built around *weighted* energy rather than raw error is that hedge-fund-style panels have wildly different target scales across instruments. Dividing by $\sum w_i y_i^2$ neutralises that scale: a 1% relative error on a low-volatility row and a 1% relative error on a high-volatility row contribute proportionally to the score. The weights w_i themselves let the organiser encode importance — and in our panel those weights are extraordinarily concentrated, so the score is effectively decided by how well the model fits a small set of high-weight rows.

3.3 Practical consequences for this project

Equation (1) is implemented once in the shared utility `cf.metrics.weighted_skill_score` and used by every later notebook for both model selection and final scoring. The function `cf.metrics.bootstrap_skill` returns 95% confidence intervals by row resampling; we report bootstrap intervals alongside every leaderboard number to avoid declaring winners that are within sampling noise of each other.

Two structural observations carry through to every subsequent modelling decision:

- **High-weight rows decide the leaderboard.** Because approximately 64% of weight sits in the top 1% of rows (Section 4.4), the score is essentially a measure of how well the model predicts those few rows. We must therefore train with sample weights, not just evaluate with them.
- **Plain RMSE will lie.** A model that improves average behaviour at the cost of a small degradation on a high-weight row can lose to a simpler model. We never report unweighted RMSE as a primary headline number.

Section 4. Dataset and Exploratory Analysis

This section reports the exploratory analysis of the panel before any modelling. The aim is twofold. First, to establish the basic statistical shape of the data so that later modelling choices can be motivated rather than guessed. Second, to surface three structural findings about the panel — multi-horizon coupling, the cumulation hypothesis, and per-horizon weight dominance — that the midway report did not document but that materially change the modelling space. All numerical results below are produced by notebook `00_data_structure.ipynb` and the corresponding figures are saved as both PNG and PDF assets.

4.1 Schema, panel scale, and train–test comparability

The training panel contains 5,337,414 rows, organised by 23 distinct `code` values, 180 distinct `sub_code` values, and 5 `sub_category` values across 4 horizons. The integer time index `ts_index` runs over the contiguous range $[1, 3601]$. The test partition continues from `ts_index` = 3602 to 4376 with 1,447,107 rows, so test is a strict temporal suffix of train. Treating the 4-tuple (`code`, `sub_code`, `sub_category`, `horizon`) as the natural series identifier yields 36,923 distinct series.

The only columns present in train but not in test are y_{target} and `weight`, as expected. Every category level seen in test exists in train, so there is no cold-start schema issue. The sub-code support narrows from 180 in train to 47 in test — a coverage shift rather than a schema mismatch.

4.2 Missing values

Missingness is computed on a 500,000-row stratified sample to keep memory in check. Of the 86 anonymised features, 48 contain at least one missing value, but the maximum missing rate across all columns is approximately 12.5%. Figure 1 shows the top fifteen offenders.




Figure 1: Top fifteen features by missingness, on a 500,000-row stratified sample of train. The maximum missing rate is approximately 12.5%, low enough that median imputation inside a fold is sufficient and there is no need to drop columns or use complex imputers.

Reading. Missing rates are bounded; within-fold median imputation suffices. No separate imputation strategy is developed.

4.3 Target distribution

The target y_{target} is centred very close to zero (median -0.0006 , IQR ± 0.05), but has extremely heavy tails. The 1% and 99% quantiles sit at approximately -82.8 and $+62.9$ respectively, while the full range runs from approximately -2202 to $+2314$.




Figure 2: Raw target distribution (left) and the same distribution clipped to $[-10, 10]$ (right). Most of the mass concentrates near zero, with rare extreme tails up to the order of ± 2200 .

Reading. The shape disqualifies plain unweighted MSE as a sensible loss. Tails are real outcomes, not corruption, so we cannot simply clip them; combined with the weighting scheme below, this is the central methodological constraint.

4.4 Weight distribution and metric concentration

The competition weights are extraordinarily skewed. The median weight is 1,699, while the maximum is 1.39×10^{13} . Figure 3 shows the weight distribution clipped at the 99th percentile (left) and the cumulative-share Lorenz-style curve (right). The top 1% of rows carry 64.2% of total weight; the top 5% carry 92.6%; the top 10% carry 98.3%.




Figure 3: Weight distribution clipped at the 99th percentile (left) and cumulative share of total weight ranked by weight (right). Approximately 64% of total weight is concentrated in the top 1% of rows; 92% in the top 5%.

Reading. Unweighted evaluation describes a different competition. The weighted skill score (Section 3) is our only headline metric, and a model that accurately predicts the top one percent

of rows has more leverage on the leaderboard than a uniformly mediocre one.

4.5 Series coverage and split feasibility

Not every series has enough history to support time-based validation. Figure 4 shows the distribution of per-series lengths. Out of 36,923 distinct series, only 1,930 (approximately 5%) cleanly cross the train/validation cutoff at $\text{ts_index} = 2880$. Many series end before the cutoff, contributing zero validation rows.




Figure 4: Distribution of per-series lengths. Most series are short, and the tail of long series carries the structural signal that classical fits can exploit.

Reading. Per-series classical fitting is feasible only on the eligible subset (cross the cutoff with at least 120 training rows). Global pooled models are not length-constrained — one reason the comparison between local classical and global ML approaches is not perfectly apples-to-apples.

4.6 Representative-series selection

For every per-series figure in this project we plot the same four series. They are selected with explicit reasons rather than at random, in order to span four distinct stress regimes: long history, large weight, high variability, and near-flat behaviour. Table 1 lists the four chosen series. They match the selection used in the midway report, preserving visual continuity.

Table 1: Four representative series chosen for per-series studies. Each series probes a different regime of the panel.

Reason	Series key (code/sub_code)	Horizon	Length
Longest history	X9BZ68VQ / OYJGNSQK / DPPUO5X2	1	212
Highest total weight	SJZP0OVU / OYJGNSQK / NQ58FVQM	25	156
Most volatile	W4S29LF4 / KL66VIS3 / PHHHVYZI	25	162
Most stable	SJZP0OVU / OYJGNSQK / NQ58FVQM	1	170

Figure 5 shows the raw y_{target} trajectory of each chosen series with the train/validation cutoff marked. The four panels look profoundly different: the most-volatile series swings between -500 and $+500$, while the most-stable series occupies a numerical range narrower than 0.001 .




Figure 5: Raw y_{target} for the four chosen representative series. The dashed vertical line marks $\text{ts_index} = 2880$, the train/validation cutoff. The four panels span the four stress regimes that any model has to handle.

Reading. A one-size-fits-all model will struggle: the most-stable series is essentially numerical noise, while the most-volatile series has obvious cyclical structure. Forecast performance across these four panels is one of the cleanest diagnostics for whether a model family generalises or only fits one regime.

4.7 Stationarity diagnostics

Classical ARIMA-family forecasting assumes stationarity, or stationarity-after-differencing. We therefore run the Augmented Dickey–Fuller test on a stratified sample of 200 eligible series. The ADF test takes the null hypothesis that the series has a unit root (non-stationary); we reject at the 5% level when the p -value is below 0.05. Figure 6 summarises the results.

Approximately 70.5% of sampled series reject the unit-root null at the 5% level. The picture is more nuanced when broken down by horizon: H1 series are more often stationary than longer-horizon cumulative series, which is consistent with the multi-horizon coupling we will document in Section 4.14.

ADF and KPSS have opposite null hypotheses, so running both gives a more honest joint verdict: ADF takes non-stationarity as the null, while KPSS takes stationarity as the null. We classify a series as *stationary* when ADF rejects and KPSS does not, as *unit root* when ADF fails to reject and KPSS rejects, and as *inconclusive* otherwise. Figure 7 reports the joint distribution.

The stationary share dominates, but the inconclusive bucket is large enough that one-size-fits-all claims about raw levels would be unsafe. We need a repair step (differencing) for the unit-root and inconclusive cases.



Figure 6: ADF results on 200 sampled eligible series. Left: p -value distribution with the 0.05 threshold marked. Centre: fraction stationary by horizon. Right: fraction stationary by code.




Figure 7: Joint ADF + KPSS verdict on 200 sampled series. Stationary 64.0%, inconclusive 29.5%, unit root 6.5%.

Figure 8 re-runs the ADF test on the first-differenced series. The stationary share rises from 70.5% on raw levels to 98.5% after differencing.

Reading. The diagnostic chain motivates ARIMA with $d = 1$ as the default for the classical study and justifies engineered differenced features for global ML models. The 98.5% post-difference stationarity also bears on the cumulation finding in Section 4.14: differencing a higher-horizon target partially inverts the rolling-sum construction back into the H1 process.

4.8 Visual stationarity — rolling statistics

Statistical tests can disagree on short or noisy series, so a rolling-window view is a useful visual check. Figure 9 shows a 24-step rolling mean and rolling standard deviation on the raw four representative series. The most-volatile series shows clear drift in level and bulging variance; the longest-history series drifts upward; the high-weight and most-stable series occupy near-flat ranges where the rolling mean barely moves.

Figure 10 repeats the same diagnostic on the first-differenced versions of the same series. The rolling means now hug zero and the rolling standard deviations flatten across the cutoff, visually confirming the statistical jump from the previous subsection.

Reading. Differenced series behave like the white-noise regime we want before fitting a differenced classical model.




Figure 8: ADF stationary share at the 5% level: levels (left, $d = 0$) versus first differences (right, $d = 1$). Differencing pushes the share to nearly 100%.

4.9 Lag structure — ACF and PACF

The autocorrelation function (ACF) shows correlation between the series and lagged copies of itself; the partial autocorrelation function (PACF) isolates each lag’s direct contribution after controlling for shorter lags. A sharp PACF cutoff suggests an autoregressive order; a sharp ACF cutoff suggests a moving-average order. Rather than relying on a single anchor series, we aggregate the absolute ACF and PACF across all 200 sampled eligible series (Figure 11). The panel-level view exposes which lags are reliably informative across the panel rather than artefacts of one trajectory.

Reading. Short-memory dynamics dominate (lags 1–3 carry most of the structure with rapid decay thereafter), and there are no clean seasonal bumps. The lag feature set for the global ML models therefore prioritises lags 1, 2, 3, 5, 10 and skips seasonal indicators.

4.10 STL decomposition

Seasonal-Trend decomposition using LOESS (STL) splits a series into trend, seasonal, and residual components. We apply STL to the longest-history anchor with a period of 24 and observe the relative amplitudes of the three components. Figure 12 shows the result.

Reading. The trend captures slow drift, but the seasonal amplitude is small relative to the residual. Combined with the absence of seasonal bumps in the panel ACF, this confirms periodic structure is not the dominant signal — modelling effort goes into short-memory dynamics, not seasonal indicators.

4.11 Feature audit

The 86 `feature_*` columns are anonymised and we cannot interpret them by name. We can, however, audit them statistically. On a 200,000-row stratified sample we compute Pearson correlation between each feature and y_{target} , then rank features by absolute correlation. Figure 13 shows the top twenty.

Reading. The single largest absolute correlation is below 0.10. No single feature predicts the target well on its own. A separate correlation check among the top twenty features further reveals strong feature–feature blocks, indicating real redundancy in the engineered set. Together




Figure 9: Rolling mean and rolling standard deviation (window = 24) on the raw four representative series. The most-volatile series shows clear drift in level and bulging variance.

these say: the signal is weak, spread across many correlated features, and best exploited by gradient-boosted trees combined with caution when reading importance plots.

4.12 Cross-series relationships

Within the same `code`, different sub-codes might share signal. If they do, global pooled models can borrow strength across related sub-codes; per-series local models cannot. We compute lagged cross-correlations between random pairs of sub-codes within the same (`code`, `horizon`) group. Figure 14 shows six representative pairs.

Reading. Clear lead/lag peaks within a code mean global models can learn a shared cross-sub-code component that per-series fits cannot. This is one of the strongest a-priori arguments for the global ML approach.

4.13 Multi-horizon coupling

The midway report treated (`code`, `sub_code`, `sub_category`, `horizon`) as the natural series identifier and fit each horizon independently. The exploratory analysis in this project shows that the four horizons of the same triple are not independent. For approximately 1.21 million distinct (`code`, `sub_code`, `sub_category`, `ts_index`) tuples, all four horizons are simultaneously present.

A second consequence of multi-horizon coupling is per-horizon target scaling. Figure 15 shows the target standard deviation and maximum row weight by horizon.

The target standard deviation rises from approximately 11.7 at H1 to 52.8 at H25, a factor of 4.5.




Figure 10: Rolling mean and rolling standard deviation (window = 24) on the first-differenced four representative series. The mean oscillates close to zero and the variance is stable, consistent with stationarity.

Reading. This finding has direct modelling consequences. A single model trained on horizons pooled into one DataFrame, without horizon-aware standardisation, will be dominated by H25 magnitudes and predict at H25 scale. Predictions at H25 scale evaluated against H1 targets will overshoot by a factor of roughly 4.5, producing catastrophically large weighted RMSE on H1 rows. This is exactly the failure mode observed in earlier ML artefacts attached to the project — the original XGBoost run reported a weighted RMSE more than two hundred times the naive baseline, consistent with a per-horizon-scale leak. The remedy is straightforward: train one model per horizon, or normalise the target by horizon before training.

4.14 Cumulation hypothesis

Given the multi-horizon coupling above, the natural hypothesis about how the four horizons relate to each other is that the target at horizon k is a cumulative aggregate of the H1 target over the next k steps:

$$y_t^{H_k} \approx \sum_{i=0}^{k-1} y_{t+i}^{H_1}. \quad (2)$$

We test Equation (2) on a stratified sample of 500 triples that have all four horizons and at least 60 H1 rows. For each triple we pivot the data so that horizons sit on columns, compute the rolling sum of the H1 series over the appropriate window, and report the Pearson correlation between $y_t^{H_k}$ and that rolling sum. Figure 16 shows the distribution of correlations and the median ratio.

Table 2 reports the summary statistics. At $k = 3$ the median Pearson correlation is 0.96




Figure 11: Panel-level mean $|ACF|$ (left) and mean $|PACF|$ (right) across 200 sampled eligible series. Persistent low-order lags dominate; higher lags decay quickly. There are no clean seasonal bumps.

and the median ratio is 1.00, indicating essentially exact aggregation. The relation is slightly noisier at $k = 10$ and $k = 25$, with median ratios of 0.95 and 0.91, but still clearly dominant.

Table 2: Cumulation hypothesis test summary on 500 triples.

k	n triples	mean corr	median corr	median ratio
3	500	0.92	0.96	1.00
10	500	0.86	0.94	0.95
25	500	0.78	0.91	0.91

Reading. Two consequences for the modelling work that follows. First, fit-H1-and-aggregate becomes a viable, possibly superior strategy to fitting per-horizon models independently: a single H1 model gives H3, H10, and H25 forecasts for free, and the aggregation is mathematically forced by the data construction. We test this in a dedicated H1-aggregation experiment in a later notebook. Second, differencing operations on $y^{H_{25}}$ partially recover the y^{H_1} process, which explains why ARIMA with $d = 1$ wins the most-volatile H25 series in the midway classical study: differencing inverts the cumulation. The midway interpretation that ARIMA was capturing a difficult pattern is partly correct; it was also undoing the cumulative construction.

4.15 Per-horizon weight dominance

The weight distribution analysis in Section 4.4 was panel-wide. Because `weight` and `horizon` are correlated, the rows that drive the weighted skill score may not be evenly distributed across horizons. Figure 17 reports the share of total panel weight per horizon and the share of weight within the top 1% of all weight mass.

The headline numbers are striking. The maximum panel weight (1.39×10^{13}) is at H3, not H1. H3 alone holds 37.4% of total panel weight and 43.6% of the top-1% weight mass; H1 holds 25.4%, H10 holds 14.3%, and H25 holds 16.8%.

Reading. A model that predicts H3 well on the small set of high-weight rows will tend to win the leaderboard regardless of how it performs on the other horizons. This was missed in the midway report. The practical consequence is that we report H3 weighted skill as a separate headline number alongside the panel-wide score, and we track H3 performance carefully when comparing models.




Figure 12: STL decomposition of the longest-history anchor with period 24. The trend component captures slow drift; the seasonal amplitude is small relative to residual standard deviation.

4.16 EDA-driven modelling implications

The EDA findings above constrain every modelling chapter that follows. The classical baselines (notebook 01) honour the chronological cutoff and use the four representative series. The global LightGBM models (notebook 02) train per horizon with sample weights, motivated by the per-horizon target scaling and H3 weight dominance. The H1-aggregation experiment (notebook 03) operationalises the cumulation hypothesis. The probabilistic forecasting work (notebook 04) addresses the heavy tails and weak per-feature signal directly through quantile regression and conformal calibration. The ensembles (notebook 05) fit weights on a meta slice that respects chronology. The final comparison (notebook 06) reports both panel-wide and H3-specific weighted skill scores with bootstrap intervals.




Figure 13: Top twenty features by absolute correlation with y_{target} , computed on a 200,000-row stratified sample. Green bars are positive correlations, red bars are negative.

Figure 14: Lagged cross-correlation between random sub-code pairs sharing the same code and horizon. A peak away from lag 0 means one sub-code leads or lags the other.




Figure 15: Per-horizon target standard deviation (left) and maximum row weight on a logarithmic scale (right). Target scale grows by a factor of approximately 4.5 from H1 to H25. Maximum weight is highest at H3.




Figure 16: Test of the cumulation hypothesis on 500 stratified triples. Left: distribution of Pearson correlations between $y_t^{H_k}$ and the rolling sum of y^{H_1} over k steps, for $k \in \{3, 10, 25\}$. Right: distribution of the per-triple median ratio $y_t^{H_k} / \sum y^{H_1}$. Higher and tighter is better.




Figure 17: Per-horizon weight breakdown. Left: each horizon's share of total panel weight. Right: share of the top 1% of weight mass by horizon. H3 dominates both views.

Section 5. Classical Baselines

This section establishes what per-series classical fitting can and cannot do on this panel. The four representative series from Section 4.6 are forecast under a single chronological protocol, scored under the weighted skill score, and compared against naive baselines that act as the visible floor. The headline finding is that classical local fits are *feasibility-limited and signal-limited*: only one of the four series produces a statistically significant improvement over the naive floor. We treat this section primarily as a scope demonstration that motivates the global pooled approach in Section 4.12 (and the modelling chapters that follow).

5.1 Setup and the four representative series

We reuse the four series chosen in Section 4.6 so that the visualisations carry over from the EDA. Two implementation choices differ from the midway report.

First, we use **adaptive 80/20 chronological splits per series** rather than the canonical fixed cutoff at `ts_index = 2880`. Under the fixed cutoff, two of the four series (highest-weight and most-stable) end up with only 7 pre-cutoff training points, which is below the minimum window required for ARMA/ARIMA estimation and forces those families to be skipped entirely. Adaptive splits preserve chronology — training rows always strictly precede validation rows — while ensuring every series has a usable training segment. The trade-off is that classical numbers in this section are not directly comparable, slice-for-slice, with the global ML numbers in later sections; we flag this explicitly when synthesising results in Section 4.15.

Second, we use **AICc-based automatic order selection** for AR, MA, ARMA, and ARIMA over a small grid of $(p, q) \in \{0, 1, 2, 3\}^2$. Hand-picking a winner from a 12-cell grid inflates the multiple-comparisons risk. AICc selects one order per series per family using a small-sample-corrected information criterion, so the model reported as “the AR result” is the one AICc would have chosen on the training data alone.

Table 3 summarises the per-series training and validation lengths under the adaptive split.

Table 3: Adaptive per-series chronological splits used for the classical study.

Series	Horizon	Train length	Val length	Validation range
Longest history	H1	169	43	2995–3037
Highest total weight	H25	124	32	3001–3032
Most volatile	H25	129	33	2879–2911
Most stable	H1	136	34	3013–3052

5.2 Model lineup

The lineup spans naive baselines (the floor), smoothing models (robust on tiny windows), and four flavours of order-selected ARIMA-family models.

We deliberately exclude SARIMA because the panel ACF in Section 4.9 shows no clean seasonal structure, so any seasonal period would be an unjustified hyperparameter. We also exclude AR/MA/ARMA/ARIMA orders larger than 3 because validation segments are 32–43 rows — larger orders cannot be reliably estimated on this much data even when they fit the training window.

What AIC and AICc do. AIC, the Akaike Information Criterion, scores a fitted model by trading off goodness of fit against parameter count:

$$\text{AIC} = -2 \log \hat{L} + 2k, \quad (3)$$

Table 4: Classical model lineup. AR / MA / ARMA / ARIMA orders are selected per series by AICc on the grid $(p, q) \in \{0, 1, 2, 3\}^2$ with the indicated d .

Family	Models	Role
Naive baselines	Zero, Naive (lag-1), Drift	Skill floor
Smoothing	Rolling Mean ($w=10$), SES, Holt	Robust on small windows
AR	order $(p, 0, 0)$, AICc-selected	Lagged-target dependence only
MA	order $(0, 0, q)$, AICc-selected	Serially correlated shocks
ARMA	order $(p, 0, q)$, AICc-selected	AR + MA in levels
ARIMA	order $(p, 1, q)$, AICc-selected	Differencing as the EDA-motivated repair

where $\hat{L}$ is the maximised likelihood of the model on the training data and k is the number of estimated parameters. Lower is better. The first term rewards models that explain the data; the second penalises models that explain it by adding parameters. Used as a model-selection criterion across non-nested specifications, AIC tends to pick the model with the best out-of-sample predictive accuracy on average, which is exactly the property we want when choosing an ARIMA order.

AICc is a small-sample correction to AIC:

$$\text{AICc} = \text{AIC} + \frac{2k(k+1)}{n-k-1}, \quad (4)$$

where n is the number of training observations. The correction term decays towards zero as n grows, so for large samples AICc and AIC agree, but for small n relative to k the AIC penalty is too lenient and the model selected by AIC tends to overfit. Burnham and Anderson recommend AICc whenever $n/k < 40$. With training segments of 124–169 rows and ARIMA orders up to $k \approx 8$ parameters, our n/k ratio sits between 15 and 21, well inside the AICc regime. Using AICc here is not a stylistic choice; it is the appropriate small-sample variant.

We use AICc to choose the order (p, q) within each family for each series, and we let it pick across families with different d values too: AR, MA, and ARMA fit the series in levels ($d = 0$); ARIMA fits the differenced series ($d = 1$). This is more principled than gating on the ADF result, because the ADF test is itself a noisy decision, whereas AICc is a continuous criterion that compares all candidate specifications on the same scale.

5.3 Evaluation protocol

Each model is fit on the per-series training segment and forecasts the full validation segment in one closed-form multi-step pass. This is the deployment-realistic protocol: no ground-truth feedback during forecasting. The headline number for every model is its weighted skill score on the validation segment under this protocol.

For ARIMA-family winners we additionally compute an *open-loop one-step* prediction at each validation time step — the model is given the true previous value at each step rather than its own previous prediction. Open-loop one-step is not a deployable protocol (the truth is unavailable at inference), but the gap between the two predictions is informative: it measures how much skill is lost when the model has to feed back its own predictions through the validation window. We render the one-step variant as a point cloud overlaid on the multi-step forecast in Figure 18, but do not promote it to a headline number. For the smoothing-family winners, the open-loop one-step prediction coincides with the multi-step forecast when the smoothing parameters are estimated on training data only, so we omit it.

We deliberately do not report unweighted RMSE or MAE as headline numbers: per the discussion in Section 3, unweighted error metrics can disagree with the weighted skill score on which model is best, because the high-weight rows dominate the weighted score.

5.4 Per-series leaderboard

Table 5 reports the best model per series under weighted skill, with 95% bootstrap confidence intervals from 500 resamples of the validation rows. The Naive (lag-1) skill is shown alongside as the floor that any reported result has to clear. Auto-selected ARIMA-family orders are shown in parentheses.

Table 5: Per-series leaderboard under the adaptive split. “Skill (95% CI)” is the weighted skill of the best model with a 500-sample bootstrap interval. “Naive floor” is the weighted skill of the lag-1 forecast on the same validation segment.

Series	Best model	Family	Skill (95% CI)	Naive floor
Longest history	AR(3, 0, 0)	AR	0.04 [0.00, 0.27]	0.00
Highest total weight	SES	Smoothing	0.47 [0.00, 0.65]	0.44
Most volatile	Rolling Mean ($w=10$)	Smoothing	0.80 [0.56, 0.91]	0.67
Most stable	Rolling Mean ($w=10$)	Smoothing	0.17 [0.00, 0.27]	0.00

The bootstrap intervals are the substantive part of this table. Three of the four series have intervals that include zero, meaning the apparent improvement over the naive floor is within sampling noise on a validation segment of 32–43 rows. The most-volatile series is the only case where the interval excludes zero: rolling-mean skill 0.80 with CI [0.56, 0.91] is a real positive result. For the highest-weight series the SES point estimate (0.47) is barely above naive (0.44); the wide CI down to 0.00 confirms that the apparent edge is fragile.

A second informative signal is which family wins. Smoothing wins three of four cases, and the only AR-family winner (longest history, AR(3, 0, 0)) has the lowest skill of the four. Differenced ARIMA, which won the corresponding case in the midway report, does not appear in this leaderboard.

5.5 Skip-status under adaptive splits

Under the adaptive 80/20 split, every model in the lineup successfully fits on every series. The skip-status table is therefore empty, which is itself the point of using adaptive splits: the same lineup under the canonical fixed cutoff at `ts_index = 2880` would have skipped every ARMA, ARIMA, $AR(p \geq 2)$, and $MA(q \geq 2)$ fit on the highest-weight and most-stable series, because their pre-cutoff training segments contain only 7 rows. Reporting fit feasibility separately from fit performance is therefore unnecessary here.

5.6 Forecast overlays

Figure 18 shows the per-series train and validation segments alongside the best model’s closed-form multi-step forecast, the Naive (lag-1) floor, and (where the best model is in the AR family) the open-loop one-step variant rendered as small dots.

The four panels have very different vertical scales (the most-volatile series swings between -600 and $+600$, the most-stable series occupies a range narrower than 0.0004), which is the same regime contrast first surfaced in Section 4.6. The forecast lines are flat or near-flat in three of the four panels because the best model in those cases is a smoothing model whose multi-step forecast is constant by construction. The most-volatile panel is the only case where the forecast line changes meaningfully across the validation segment, and it is also the only case where the model produces a substantively positive skill above naive. The open-loop one-step dots in the longest-history panel are visibly closer to the validation truth than the multi-step line — the visible gap is the value of feedback at inference for that series and that model.




Figure 18: Forecast overlays for the four representative series under the adaptive split. Black: training segment. Orange: validation truth. Blue solid: best model under closed-form multi-step. Blue dots: open-loop one-step prediction (only shown for AR-family winners). Grey dashed: Naive (lag-1) floor. Vertical dashed line: per-series train/validation cutoff. Skill annotations are weighted skill scores against the validation segment.

5.7 What this study can and cannot say

Three points to carry forward:

1. **Local classical fits are bounded by training length and target structure, not by model sophistication.** Auto-selecting orders, expanding the grid, or substituting fancier classical variants would not move the headline numbers in any meaningful way. The cap on classical performance here is the data, not the methodology.
2. **Smoothing wins on every series except the longest-history case, and even there the AR winner has skill 0.04.** The midway report attributed the volatile-series win to ARIMA $d=1$ “capturing differencing structure”. Under the adaptive split with auto-selected orders, the simpler Rolling Mean baseline wins instead, with the strongest statistical signal among all four series. This suggests that on this panel, robust low-complexity smoothing is at least as defensible as differenced ARIMA — and is the safer reporting choice given the bootstrap intervals.
3. **Per-series fitting throws away the global panel structure.** The 86 anonymised features and the cross-series correlations documented in Section 4.12 are entirely unused here. The next chapter lifts both restrictions by training globally with sample weights, and the per-horizon weight-dominance finding in Section 4.15 guides where the additional capacity should be spent.

This section is therefore a scope demonstration: it establishes the upper bound on what local classical fitting can achieve on this panel, with proper bootstrap intervals so that the bound is statistically defensible. The modelling work proper begins in the next chapter.